Kcits970 Problem Solving Handbook

When I learn algorithms/techniques, I often forget their correctness proofs. This handbook contains all algorithms that I have learned, but forgotten the proofs for. Please note that the proofs included here are not intended for complete rigor, but to provide a convincing, memorable argument.

1 Longest Increasing Subsequence

Reference: https://cp-algorithms.com/dynamic_programming/longest_increasing_subsequence.html
Implementation: ./code/arr/lis.cpp

---

1. Let A be the given sequence of length n.

2. Let $f = (A_1, \underbrace{\infty, \infty, \infty, \cdots, \infty}_{n-1})$.

3. For each i from 2 to n, binary search the smallest l such that $a_i \leq f_l$ and update f_l to A_i.

4. The largest l such that $f_l < \infty$ is the length of the longest increasing subsequence of A.

Define $g_{i,l}$ as the smallest terminal element across all increasing subsequences of length l within the ith prefix of A. We show that in the above algorithm, f_l is equal to $g_{i,l}$ after the ith iteration. (Then the rest of the correctness should be trivial.) Notice the following recurrence relation of g.

$$g_{i,l} = \begin{cases} A_i & (g_{i-1,l-1} < A_i \leq g_{i-1,l}) \\ g_{i-1,l} & \text{(otherwise)} \end{cases}$$

Additionally, for all i, $g_{i,l}$ must be strictly increasing in respect to l. This is because if $g_{i,l-1} \geq g_{i,l}$, we can obtain an increasing subsequence of length $l-1$ with a terminal element smaller than $g_{i,l-1}$. Therefore, for all A_i, there exists exactly one l such that $g_{i-1,l-1} < A_i \leq g_{i-1,l}$, and the recurrence relation of g matches exactly the operations being done on f. Since the initial state of f is equal to $g_{1,l}$ (trivially), the principle of induction tells us that f after the ith iteration equals $g_{i,l}$.

2 Strongly Connected Component

Reference: https://cp-algorithms.com/graph/strongly-connected-components.html
Implementation: ./code/graph/scc.cpp

---

1. Given a directed graph G, construct its transpose H.

2. Sort the vertices of G in the order of dfs exit time.

3. Find an *unvisited* vertex u with the greatest exit time, and collect all *unvisited* vertices reachable from u in H.

4. Repeat the above until all vertices are *visited*.

Consider two different strongly connected components S and T in G. If there exists an edge from S to T in the condensation graph C, the dfs exit time of S must be greater than that of

T. This can be proven by splitting the dfs entry order into two cases. (The rest is kind of trivial, so it is omitted here.)

$$\begin{cases} 1. \ S \text{ is visited first.} \cdots \\ 2. \ T \text{ is visited first.} \cdots \end{cases}$$

Therefore, the strongly connected component containing u (the vertex with the greatest exit time) must not have any incoming edges from other strongly connected components in G. Hence, by collecting all vertices reachable from u in the transpose, we end up with a strongly connected component. (This is valid because the set of all strongly connected components of G is equal to that of H.) Thus, repeating until exhaustion must correctly find all strongly connected components.

3   Cross Product of 2D Vectors

> 1. Let $\mathbf{u} = (a, b)$, $\mathbf{v} = (c, d)$, and θ be the angle from $\mathbf{u}$ to $\mathbf{v}$ measured in counterclockwise direction. $(0 \leq \theta < 2\pi)$
>
> 2. The signed area of the parallelogram formed by $\mathbf{u}$ and $\mathbf{v}$ is equal to $ad - bc$, and is notated as $\mathbf{u} \times \mathbf{v}$.

It suffices to show that $ad - bc = |\mathbf{u}||\mathbf{v}| \sin \theta$. First, we use the rotation matrix to find a different expression for $\mathbf{v}$.

$$\mathbf{v} = \sqrt{\frac{c^2 + d^2}{a^2 + b^2}} \begin{pmatrix} \cos \theta & -\sin \theta \\ \sin \theta & \cos \theta \end{pmatrix} \begin{pmatrix} a \\ b \end{pmatrix} = \sqrt{\frac{c^2 + d^2}{a^2 + b^2}} \begin{pmatrix} a \cos \theta - b \sin \theta \\ a \sin \theta + b \cos \theta \end{pmatrix}$$

Then we calculate $ad - bc$ using the alternate expression.

$$\begin{aligned} ad - bc &= \sqrt{\frac{c^2 + d^2}{a^2 + b^2}} (a(a \sin \theta + b \cos \theta) - b(a \cos \theta - b \sin \theta)) \\ &= \sqrt{\frac{c^2 + d^2}{a^2 + b^2}} (a^2 \sin \theta + b^2 \sin \theta) \\ &= \sqrt{a^2 + b^2} \sqrt{c^2 + d^2} \sin \theta \\ &= |\mathbf{u}||\mathbf{v}| \sin \theta \end{aligned}$$

This fact implies that $\mathbf{v}$ is oriented counterclockwise to $\mathbf{u}$ if $\mathbf{u} \times \mathbf{v}$ is positive, and clockwise if $\mathbf{u} \times \mathbf{v}$ is negative. The cross product of zero indicates that the two vectors are parallel to each other.

4   Orientation of 3 Points

This algorithm is also known as *CCW*.

> 1. Let A, B, C be arbitrary *distinct* points on a plane.
>
> 2. The orientation of A, B, C is counterclockwise if $\overrightarrow{AB} \times \overrightarrow{BC} > 0$, and clockwise if $\overrightarrow{AB} \times \overrightarrow{BC} < 0$.
>
> 3. If the cross product is zero, then the three points are on a line.

This is a direct derivation from the cross product of 2D vectors.

5 Line Segment Intersection Check

1. Let A, B, C, D be arbitrary *distinct* points on a plane.

2. If A, B, C and A, B, D are oriented in the same direction, then $\overline{AB}$ and $\overline{CD}$ do not intersect.

3. The same argument applies to the orientations of C, D, A and C, D, B.

4. If the four points are on a line, then the two segments intersect if and only if one of the following conditions hold.

   - The rightmost endpoint of $\overline{AB}$ is to the right of the leftmost endpoint of $\overline{CD}$.
   - The leftmost endpoint of $\overline{AB}$ is to the left of the rightmost endpoint of $\overline{CD}$.
   - The uppermost endpoint of $\overline{AB}$ is above the lowermost endpoint of $\overline{CD}$.
   - The lowermost endpoint of $\overline{AB}$ is below the uppermost endpoint of $\overline{CD}$.

The correctness of this algorithm can be shown visually, but I don't really want to write a dedicated TikZ diagram for it. Also, it should be noted that when implementing the algorithm, the usage of $<, \leq, >, \geq$ may depend on how the problem sets the condition for an *intersection*.

6 Diameter of a Tree

Reference: https://codeforces.com/blog/entry/101271
Implementation: `./code/tree/diam.cpp`

If the weights can be negative, then this algorithm does not work in the general case.

1. Let T be a given tree with non-negative weights, and p be an arbitrary node on T.

2. Let a be any farthest node from p.

3. Let b be any farthest node from a.

4. The path from a to b is one of the diameters of T.

Let P denote the path from a to b. For each vertex v, define r_v as the first *reachable* vertex on P from v. First, we show that for every vertex v, $\text{dist}(r_v, v) \leq \text{dist}(r_v, b)$. By definition of b, $\text{dist}(a, v) \leq \text{dist}(a, b)$. Subtracting $\text{dist}(a, r_v)$ on both sides immediately yields $\text{dist}(r_v, v) \leq \text{dist}(r_v, b)$.

Next, we proceed to show $\text{dist}(r_v, v) \leq \text{dist}(r_v, a)$ for every vertex v. Given an arbitrary v, the position of r_v splits into two cases.

$$\begin{cases} r_v \text{ exists on the path from } a \ (\textit{inclusive}) \text{ to } r_p \ (\textit{exclusive}). \\ r_v \text{ exists on the path from } r_p \ (\textit{inclusive}) \text{ to } b \ (\textit{inclusive}). \end{cases}$$

In the former case, $\text{dist}(r_v, v) \leq \text{dist}(r_v, a)$ must hold, because otherwise it contradicts the fact that a is farthest from p. In the latter case, $\text{dist}(r_v, b) \leq \text{dist}(r_v, a)$, because otherwise b is farther away from p than a. Combining this with the previously derived inequality $\text{dist}(r_v, v) \leq \text{dist}(r_v, b)$ returns $\text{dist}(r_v, v) \leq \text{dist}(r_v, a)$.

Finally, we show that $\text{dist}(u, v) \leq \text{dist}(a, b)$ for every pair of vertices u and v. Without loss of generality, assume that r_u exists on the path from a to r_v.

$$\begin{aligned} \text{dist}(u,v) &\leq\ \text{dist}(u,r_u) + \text{dist}(r_u,r_v) + \text{dist}(r_v,v) \\ &\leq\ \text{dist}(a,r_u) + \text{dist}(r_u,r_v) + \text{dist}(r_v,b) \\ &=\ \text{dist}(a,b) \end{aligned}$$

We have proven that $\text{dist}(a,b)$ is greater than or equal to every $\text{dist}(u,v)$; hence the algorithm correctly finds the diameter of T.